

Q1: Explain how AI-driven code generation tools (e.g., GitHub Copilot) reduce development time. What are their limitations?

How AI-Driven Code Generation Tools Reduce Development Time

AI tools like GitHub Copilot help accelerate development by:

- **Auto-generating boilerplate code:** Saves time on repetitive tasks like setting up functions, classes, or standard patterns.
- **Smart code suggestions:** Offers context-aware completions based on the current file, project, and even docstrings.
- **Built-in knowledge:** Acts like a coding assistant with an encyclopedia of APIs, libraries, and best practices.
- **Real-time error spotting:** Identifies common issues or syntax mistakes before execution.

Together, these features streamline workflows, reduce context-switching, and boost overall productivity—especially for routine or well-defined tasks.

Limitations of AI Code Generation Tools

- **Shallow understanding:** AI may generate logically incorrect or inefficient code, especially in complex scenarios.
- **Security risks:** Suggested code might include unsafe patterns or resemble copyrighted material.
- **Overreliance:** Developers might skip learning underlying concepts, impacting long-term growth.
- **Context limitations:** Performance may drop in large or unfamiliar codebases where nuance matters.

Q2: Compare supervised and unsupervised learning in the context of automated bug detection.

FEATURE	SUPERVISED LEARNING	UNSUPERVISED LEARNING
Data Requirement	Requires labeled data (e.g., code marked as buggy or clean)	Uses unlabeled data (just raw code or system logs)
Goal	Learn to classify known bug types	Discover unknown or unusual behaviors (anomalies)
Example Use Case	Training a model to detect null pointer exceptions	Identifying abnormal memory usage patterns
Strengths	High accuracy on known issues; easy performance tracking	Can detect new, rare, or subtle bugs
Limitations	Needs large, high-quality	Results can be vague or prone

	labeled datasets	to false positives
Output	Explicit bug classification (e.g., “divide-by-zero”)	Alerts about anomalies without specific labels

Supervised learning: is great when you have a rich history of bugs to learn from and want precise detection.

Unsupervised learning: shines when you want to uncover hidden problems, patterns, or emerging vulnerabilities.

Q3: Why is bias mitigation critical when using AI for user experience personalization?

Ensures fairness: AI systems learn from historical data, which may reflect real-world inequalities or stereotypes. If left unchecked, they can amplify those biases—showing certain products, features, or content only to specific groups and excluding others.

Improves personalization quality: Bias can skew predictions, leading to irrelevant or even offensive recommendations. A more neutral and inclusive model better captures diverse user preferences.

Promotes inclusivity: Everyone deserves to feel seen and represented. Bias-free personalization ensures users from different backgrounds, cultures, and identities all get engaging, relevant experiences.

Protects brand trust: Discriminatory or skewed results can damage a company’s reputation. Proactively addressing bias helps avoid public backlash and supports responsible innovation.

Supports compliance: Many regions now have legal frameworks around algorithmic fairness, especially in sectors like finance, hiring, and media. Ethical personalization keeps products on the right side of regulations.

How does AIOps improve software deployment efficiency? Provide two examples.

AIOps is an Acronym for Artificial Intelligence for IT operations that streamlines and strengthens software deployment by using machine learning and big data analytics to automate and optimize workflows.

How AIOps Enhances Deployment Efficiency

- **Intelligent automation:** AIOps reduces manual intervention by predicting issues, automatically resolving repetitive tasks, and adjusting deployment parameters on the fly.
- **Real-time monitoring and insights:** It continuously analyzes telemetry from logs, metrics, and traces to detect anomalies early—before they impact users.
- **Faster root cause analysis:** Instead of sifting through endless logs, AIOps pinpoints the likely source of errors using pattern recognition and correlation.

Example 1: Predictive Rollback with Anomaly Detection

Imagine a company like an e-commerce platform rolling out a new payment gateway. With AIOps in place:

- It monitors traffic, latency, and user behavior in real time.
- If it detects an unusual drop in successful transactions immediately after deployment, it automatically rolls back the update to maintain availability.
- Engineers receive detailed diagnostics and visualizations of what triggered the rollback—saving hours of manual debugging.

Example 2: Intelligent Resource Allocation During CI/CD

In a continuous deployment pipeline, infrastructure demands spike during builds, testing, and deployment:

- AIOps analyzes historical usage patterns and forecasts upcoming demand.
- It auto-scales environments or shifts workloads to underutilized clusters—avoiding bottlenecks and optimizing cloud spend.
- Developers experience faster builds and fewer pipeline delays without manual tuning.

Ethical Reflection

Potential Biases in the Dataset

Demographic Imbalance: If the dataset overrepresents certain age groups, ethnicities, or genders, the model may generalize poorly to underrepresented populations.

Labeling Bias: If issue priority is derived from subjective or historical decisions, it may reflect past human biases.

Feature Bias: Some features may correlate with protected attributes (e.g., socioeconomic status), introducing indirect bias.